

Genre Comparison through Music Signal Processing

Jeremy Grifski

Abstract— Genre is a popular way to classify styles of music, and people may argue there is merit in genre classification. But, are there any tangible differences in music styles that can be discovered through data visualization? In this project, I attempt to answer that question by applying a handful of signal processing methods to a set of music files and visualizing the results relative to genre. All code for the project can be found at <https://github.com/jrg94/CSE5539> and the live visualizations can be viewed at <https://jrg94.github.io/CSE5539/>.

Index Terms—Signal Processing, Python, JavaScript, D3

1 INTRODUCTION

As mentioned in the proposal, I am a first-year PhD student with a limited research background. My interests are mainly in music, education, gaming, and data visualization, so it is probably no surprise that I chose to do a project with two of those interests in mind: music and data visualization.

My goal with this project was to explore different genres of music through data visualization while dipping my toes in some basic signal processing. In particular, I chose to analyze a personal collection of songs in two stages: **data collection** and **data visualization**.

In terms of data collection, I spent some time learning how to parse a common file format, M4A, and interpret that data. As for data visualization, I spent some time trying to address tasks such as:

- How frequent are the various genres in my personal library?
- How does music loudness vary by genre?
- How has music duration varied over time?

In the following sections, I will further explore the processes and results of data collection and visualization.

2 DATA COLLECTION

In order to visualize the differences in music genre, I had to first collect some data. Fortunately, I had quite the collection of songs from iTunes purchases over the years. In particular, I had 4,194 songs spread across 983 folders.

From there, I had two challenges: choosing which file format to process and selecting what data to collect from those files. Of the 4,194 songs, roughly 1,800 of them were in the M4A format, so I chose to parse that format.

In terms of data collection, I decided to mine several fields from the M4A files such as the fields found in Table 1. Many of these fields were available directly in the M4A file. However, some of the fields had to be calculated using information provided in the file. For example, duration d was not provided directly. It had to be computed from the sample size n , number of channels c , sample rate r , and audio size a , or more formally:

$$d = a / (c * n * r)$$

Here, I am assuming that audio size and sample size are both in bytes. As a result, a song with two channels (stereo), a 44,100 Hz

- Jeremy Grifski is an OSU PhD Student. E-mail: grifski.1@osu.edu.

Manuscript received xx xxx. 201x; accepted xx xxx. 201x. Date of Publication xx xxx. 201x; date of current version xx xxx. 201x. For information on obtaining reprints of this article, please send e-mail to: reprints@ieee.org. Digital Object Identifier: xx.xxx/TVCG.201x.xxxxxxx

Table 1. Data Fields

field	type
Title	String
Genre	String
Artist	String
Album	String
Purchase Date	Date
Release Date	Date
Track Number	Integer
Sample Rate (Hz)	Integer
Sample Size (bits)	Integer
Duration (HH:MM:SS)	String
Number of Channels	Integer
Average dBFS	Float
Max dBFS	Float
RMS	Integer
Max Amplitude	Integer

sample rate, a two-byte sample size, and a 176,400-byte audio size would contain just one second of audio.

Likewise, I also had to compute average decibels relative to full scale (dBFS), maximum dBFS, root-mean-square (RMS), and maximum amplitude which are all parameters related to audio loudness [8]. Fortunately, I did not compute these parameters myself as they were provided through the Python pydub library [9].

In the following sections, I'll discuss how I parsed the M4A files, and how I aggregated that data for visualization.

2.1 File Parsing

The first major challenge in this project was finding a way to collect information from the M4A file format. Unfortunately, I was unable to find a tool which could easily glean all the data I wanted from an M4A file. As a result, I decided to write my own file parsing code in Python.

With the aid of the Quicktime documentation, I was able to write a Python script which could roughly interpret the various atoms of an M4A file. In particular, I wrote enough code to parse all of the atoms seen in Table 2.

An atom is a section of an M4A file that stores information based on some accepted standard [1]. Some atoms contain atoms, so parsing can be challenging. Fortunately, Python makes it relatively easy to read and parse byte streams. The challenge was properly mapping each atom based on the standard.

In total, the M4A parsing code is roughly 800 lines of code, yet it is still incomplete. In particular, I was unable to parse the actual audio data as M4A files are in many cases compressed. To account for this setback, I leveraged a Python library called pydub which uses FFmpeg to parse audio files of many formats including M4A [9].

Table 2. All Parsed Atoms

atom	description
moov	the movie atom
trak	the trak atom
mdia	the media atom
minf	the media information atom
stbl	the sample table atom
udta	the user data atom
dinf	the data information atom
pinf	-
sch	-
ftyp	the file type atom
mvhd	the movie header atom
hdlr	the handler atom
mdhd	the media header atom
dref	the data reference atom
smhd	the sound media information header atom
stco	the chunk offset table atom
stsc	the sample-to-chunk table atom
stsd	the standard description atom
esds	the elementary stream descriptor atom
stsz	the sample size table atom
stts	the time-to-sample table atom
tkhd	the track header atom
meta	the meta data atom
ilst	the item list atom
cpil	the compilation atom
gnre	the content genre (i.e. rock) atom
rtng	the content rating (i.e. explicit) atom
stik	the media type (i.e. movie) atom
mdat	the media data atom
frma	the data format atom

2.2 Data Aggregation

Being able to parse an M4A file and being able to use that data for visualization are two separate challenges. As a result, I had to find a way to aggregate the parsed data in an easy-to-read format. With the target environment being D3 and JavaScript, I chose to format all the data as JSON files. JSON is easy to load in JavaScript and even easier to change as needed. An example JSON might look something like the following:

```
{
  "path": "E:\\Plex\\Music
\\A Loss for Words\\The Kids Can't Lose
\\01 Stamp of Approval.m4a",
  "genre": "Alternative",
  "title": "Stamp of Approval",
  "artist": "A Loss for Words",
  "album": "The Kids Can't Lose",
  "track_number": 1,
  "total_tracks": 11,
  "content_rating": null,
  "owner": "Jeremy Grifski",
  "release_date": "2009-05-12T07:00:00Z",
  "purchase_date": "2014-02-10 16:44:15",
  "sample_rate": 44100,
  "sample_size": 16,
  "length": "00:03:06",
  "number_of_channels": 2,
  "dBFS": -11.080573225629289,
  "max_dBFS": 0.0,
  "rms": 599654648,
  "max_amplitude": 2147483648
}
```

With a JSON like this, it's very easy to retrieve information about a song like its genre or title. In addition, the JSON can be easily expanded to include information like the occurrences of certain pitches.

In order to generate the JSON, I created two Python classes: MusicFile and MusicFileSet. A MusicFile provides music file data storage as seen in the JSON above. In addition, the MusicFile class provides extended functionality like the ability to convert an M4A to a WAV file and the ability to persist the parsed atom data to JSON [6].

Each MusicFile can then be stored and manipulated in an instance of the MusicFileSet class. The MusicFileSet class is essentially a list wrapper, but it provides several key functionalities such as [7]:

- the ability to load all M4A files from a directory recursively
- the ability to filter a MusicFileSet by genre or size
- the ability to dump the entire set to JSON

Naturally, I leveraged the MusicFileSet class to build a large collection of data files which encompass various sample sizes and genres. In total, I was able to generate 33 data sets.

3 DATA VISUALIZATION

The goal of this project was to learn some signal processing while trying to find interesting trends in genres of music. To do so, I felt it was appropriate to try to plot some of the data I collected in D3.

3.1 Genre Histogram

To begin, I felt it was important to get an idea of what kind of data I was using. Since I was planning on working with genre data, I decided to create a histogram of the various genres. Needless to say, the results were skewed toward my favorite genres as seen in Fig. 1.




Fig. 1. A genre histogram from [5] of all the files in the dataset.

The visualization in Fig. 1 was rendered from Table 3 which contains 17 distinct genres including a placeholder for songs that have no documented genre. For convenience, the table was sorted in the same order as the plot. As a result, it's clear which genres dominate the collection: alternative and rock.

To encode genres in later plots, I chose to use color. Unfortunately, D3 no longer supports color palettes larger than 12 distinct colors, so I was forced to filter the data set by the 12 most frequent genres in the dataset. As a result, there are about 12 data points missing from Fig. 2

Table 3. Genre Totals

Genre	Count
Alternative	448
Rock	368
None	353
Pop	207
Jazz	130
Soundtrack	107
Country	88
Comedy	44
Metal	42
Punk	31
Rap	20
Classical	12
Electronic	7
Vocal	6
Dance	3
Hip-Hop	2
Instrumental	1
Sum	1869

and Fig. 3. In particular, the vocal, dance, hip-hop, and instrumental genres have not been included. Instead, the contain only the genres contained in Table 4.

Table 4. Filtered Genre Totals

Genre	Count
Alternative	448
Rock	368
Pop	207
Jazz	130
Soundtrack	107
Country	88
Comedy	44
Metal	42
Punk	31
Rap	20
Classical	12
Electronic	7
Sum	1504

Since songs lacking genre data provide no useful insight in a genre comparative study, I chose to exclude all files with None tags for genre. In total, I had about 1,500 music files for visualization purposes.

3.2 Duration Versus Release Date

With a proper head count of genres and files, I proceeded to render a visualization of duration by release date which can be seen in Fig. 2. After all, I was interested to see if song durations have changed over time.

Right away, it's clear which genres tend to have the longest durations. For example, Metal and Classical music look like outliers with their durations almost three times as long as what appears to be the average duration. I also find it interesting that duration appears to be decreasing over time. That said, I suppose that makes sense. Three-minute pop songs are a lot more common today than 8-minute jazz songs akin to Maynard Ferguson's rendition of Hey Jude.

In addition, I'm interested in seeing how the genres cluster. Clearly, some genres are more popular today than in the past, and many of them seem to follow the same formula in terms of duration. Apparently, consumers have chosen some sort of duration sweet spot which hovers around 200 seconds.

3.3 dBFS Versus Release Date

Since I was on the subject of data over time, I thought it would be interesting to look at loudness over time using decibels at full scale




Fig. 2. A plot of duration versus release date from [4] where colors correspond to genres.

(dBFS) and release date. The ideas here being that songs may have changed in average loudness over time. Naturally, the dBFS versus release data plot can be seen in Fig. 3.

After looking at Fig. 3, it's tough to make a conclusion on any trends in loudness over time. However, there are a few key features to notice. For one, loudness appears to be trending up slightly. However, the loudness is definitely more spread now than it used to be. In general, it seems songs used to hover between -12 and -22 dBFS. Now, songs range anywhere from -8 to -26 dBFS.

In terms of genre, there are a lot of interesting trends. For example, Rock is very clearly trending up in loudness. In the 70s, rock could be seen between -14 and -24 dBFS. Now, rock can be seen at -14 in its quietest. Likewise, genres like alternative and metal appear to be relatively loud on average. Meanwhile, the soundtrack genre tends to be very quiet.

3.4 Average dBFS Versus Genre

Having seen an overview of dBFS over time, I had to take the analysis one step further by asking: what genres are loudest on average? To do that, I ran an analysis which captured the average dBFS of each genre and plotted on a bar graph (Fig. 4) similar to the genre histogram.

In Fig. 4, I noticed right away that dance is the loudest genre by a considerable margin, and I suppose that makes sense. If the target audience of dance music is the club, it make sense for dance songs to be quite loud. However, it's important to notice that dance has only 3 samples, so take that graph with a grain of salt.

Closely behind dance is punk and hip-hop. Again, hip-hop only has two samples representing that average, so that genre can be safely ignored. Meanwhile, punk has around 31 samples which may all be from one artist, so it's tough to really come to any conclusions in that regard.

That said, from a biased point-of-view, I think this curve is fairly accurate. I would expect rougher genres like metal, punk, and dance to be louder on average than genres that have more ebb and flow like jazz and instrumental. Also, I'd expect genres with less simultaneous sounds like vocal to be quieter on average.

Perhaps most surprisingly, classical comes in at the quietest genre from its twelve sample average. My experience with live classical




Fig. 3. A dBFS versus release date from [2] where colors correspond to genres.

music is that it can be very quiet at times and very loud at other times. Perhaps that contrast is skewed toward the lower end for more effect. I would be interested in digging into this more in the future.

4 CHALLENGES

As it turns out, the scope of the initial proposal was much too broad. Many of the signal processing approaches proposed were outside of my capabilities within a semester. Given more time, I would have liked to explore audio features beyond loudness such as pitch and onset.

In addition, it turns out that reading and parsing existing audio formats is really challenging. Not only are the formats themselves not well documented, but the files are general fairly large. As a result, processing can take a lot of time and resources. Retrieving data for visualization took many hours of periodic persistence to ensure nothing was lost.

Finally, overall software design was challenging. From parsing and data collection to data aggregation and visualization, there were a lot of steps in the pipeline. In particular, I ended up leveraging two separate scripting languages, Python and JavaScript, as well as an intermediate data representation, JSON. In total, I have written over 1600 lines of code for this project not including the report.

5 CONCLUSION

After exploring M4A file parsing, audio signal processing, and genre data visualization, I have concluded that analyzing music is a very difficult task. While I had plenty of data, parsing common file formats, analyzing audio levels, and generating compelling visualizations are not easy tasks.

That said, I learned a lot about the relationships between different genres of music. For example, it's hardly surprising that certain genres have longer songs on average (i.e. Jazz, Classical, and Metal). However, I was surprised to see that songs have gotten slightly shorter over time.

In the future, I'd like to explore music genres with an even larger data set as to avoid introducing any personal biases.




Fig. 4. A average dBFS histogram from [3] of all the files in the dataset.

REFERENCES

- [1] Apple. Movie Atoms. <https://developer.apple.com/library/archive/documentation/QuickTime/QTFF/QTFFChap2/qtff2.html>, 2016 (accessed 09-April-2019).
- [2] J. Grifski. dBFS Vs. Release Date Plot - CSE 5539. <https://jrg94.github.io/CSE5539/term-project/musiviz/dbfs-vs-release-date.html>, 2019 (accessed 09-April-2019).
- [3] J. Grifski. dBFS Vs. Release Date Plot - CSE 5539. <https://jrg94.github.io/CSE5539/term-project/musiviz/average-dbf-vs-genre.html>, 2019 (accessed 09-April-2019).
- [4] J. Grifski. Duration Vs. Release Date Plot - CSE 5539. <https://jrg94.github.io/CSE5539/term-project/musiviz/duration-vs-release-date.html>, 2019 (accessed 09-April-2019).
- [5] J. Grifski. Genre Histogram - CSE 5539. <https://jrg94.github.io/CSE5539/term-project/musiviz/genre-histogram.html>, 2019 (accessed 09-April-2019).
- [6] J. Grifski. MusicFile. https://github.com/jrg94/CSE5539/blob/master/term-project/musiparse/music_file.py, 2019 (accessed 09-April-2019).
- [7] J. Grifski. MusicFileSet. https://github.com/jrg94/CSE5539/blob/master/term-project/musiparse/music_file_set.py, 2019 (accessed 09-April-2019).
- [8] J. Price. Understanding dB. <http://www.jimprice.com/prosound/db.htm>, 2007 (accessed 09-April-2019).
- [9] J. Robert, V. Srinivasan, and T. Fiers. API Documentation. <https://github.com/jiaaro/pydub/blob/master/API.markdown>, 2018 (accessed 09-April-2019).